

Visibility Friction in the Hong–Kangxi Discourse on Xiaohongshu

Yidan Xia, Political Science

Project: Xiaohongshu (XHS) Hong–Kangxi posts visibility study.

Prepared for: Imagined stakeholder is the proposal review committee, including researchers familiar with computational text analysis. They will decide whether the project’s measurement architecture and inferential framing are sound enough to be developed into fuller thesis analysis. The reader is assumed to be comfortable with regression and intercoder reliability but not with XHS-specific slang or the Hong–Kangxi rumor.

Executive summary

This memo asks whether public-search visibility friction inside one Xiaohongshu (XHS) Hong–Kangxi discourse event is associated with narrative direction, rhetorical mode, carrier material, or platform legibility. The practical problem for the thesis is not whether the platform “deleted” a topic in general, but whether the measurement architecture can support a narrower claim about differential visibility inside a single historical-nationalist discourse. The analytic corpus contains 1,594 unique event posts from a November 2025 keyword crawl. The outcome is **search_nonretrievable**: whether a collected post could still be recovered through public XHS search during the April 2026 verification check. Because the audit shows that this label mixes hard disappearance and search invisibility, the memo treats it as *visibility friction*, not confirmed deletion or censor intent.

The main result is descriptive. After excluding verification errors and mode/content-uncodable posts, the adjusted logit sample is $N = 1,463$. The clearest associations are feminist-lineage direction (OR = 2.06, 95% CI [1.22, 3.46], AME = +16.5 percentage points), Red Chamber suoyin carrier (OR = 2.30, 95% CI [1.64, 3.21], AME = +19.2 points), and medium OCR legibility relative to high-caption posts (OR = 1.37, 95% CI [1.05, 1.79], AME = +7.3 points). The OCR result is best interpreted as a format/searchability mechanism, not as evidence that OCR posts were deleted. Engagement is also heavily right-skewed, and very low engagement is associated with higher search non-retrievability; I therefore treat engagement as a visibility mechanism and sensitivity concern rather than a clean pre-treatment control. Anti-Manchu/Qing direction is positive but not statistically stable in the main model; anti-West is too rare to interpret beyond low-power evidence. New engagement and OCR/legibility sensitivity checks show that Red Chamber suoyin remains robust after excluding low-engagement or OCR-dependent posts, while feminist-lineage remains positive but becomes less precise in the strict high-caption-only sample. The thesis should therefore frame the quantitative finding as a within-event visibility association and reserve stronger claims about mechanisms for future verification, audit, and qualitative follow-up.

Decisions for the reader

- Use **search_nonretrievable**, but name it “visibility friction.” Do not call it confirmed deletion unless later verification separates search invisibility, author deletion, and platform removal.

- Keep direction as the central theoretical construct, but report the current anti-Manchu/Qing result as unresolved rather than confirmed. The supported direction result is feminist-lineage visibility friction.
- Treat carrier variables as material controls and sensitivity checks. Red Chamber suoyin is empirically strong, but the coefficient should not be converted into a claim that “literary content” is targeted.
- Treat medium OCR legibility as a searchability/observability mechanism. Report OCR-restricted sensitivity checks so that format-driven non-retrievability is not confused with deletion or content moderation.
- Treat engagement as a visibility mechanism and sensitivity issue, not as a clean pre-treatment control. Report models with and without engagement, categorical engagement-bin controls, and low-engagement exclusions.
- For thesis-stage causal or mechanism claims, add repeated verification, stable account identifiers, and a random audit of non-adjudicated posts.

1 Problem, scope, and research question

The Hong–Kangxi rumor claimed that the Kangxi Emperor was the secret son of Empress Dowager Xiaozhuang and Hong Chengchou, a Ming general who surrendered to the Qing. The historical claim is not the object of this memo. The research problem is that the rumor reactivated several politically meaningful ways of narrating Chinese nationhood: anti-Manchu/Qing rupture, anti-Western pseudo-history, and feminist-lineage critique. These directions touch different premises in the state-led national story: the multi-ethnic nation framework (Mullaney, 2011; Brubaker, 2004; Wimmer, 2013), the national-humiliation chronology that centers nineteenth-century foreign aggression (Wang, 2012), Han-centered Ming/Qing boundary narratives (Fitzgerald, 1996; Carrico, 2017), and gendered or patrilineal assumptions about lineage continuity (Wang, 2017).

Chinese information control is known to be strategic, delegated, and frictional rather than reducible to universal deletion (King, Pan, and Roberts, 2013; Lorentzen, 2014; Roberts, 2018; Sun and Zhao, 2022). The missing piece here is intra-domain variation: **within the same Hong–Kangxi event corpus, which post features are associated with reduced public-search visibility?** I answer three operational questions. RQ1 describes the corpus by rhetorical mode, narrative direction, carrier material, and platform legibility. RQ2 tests whether anti-Manchu/Qing, anti-West, and feminist-lineage directions differ in visibility. RQ3 tests whether mode, carrier, or legibility help explain visibility beyond direction.

2 Data, outcome, and measurement design

The raw November 2025 crawl contains 2,797 keyword hits from 康熙瓜, 洪承畴, and 伪史. After removing 107 unrelated keyword-noise posts, the cleaned crawl contains 2,690 rows: 593 standalone Western pseudo-history posts, 503 pure historical-analysis posts, and 1,594 Hong–Kangxi event posts. The 1,594 event posts are the analytic corpus. Each record includes title, caption/body, OCR text where available, post type, keyword source, engagement metadata at collection, posting date, and April 2026 search-verification status.

The dependent variable is `search_nonretrievable`: 1 if verification search does not return the original post and 0 if it does. A random audit of 100 posts flagged likely deleted found 32 gone from

an identifiable profile, 27 still present on an author profile but invisible through public search, and 41 author-unverifiable cases. This is why the outcome is interpreted as public-search visibility friction rather than hard deletion. Appendix D reports the audit and reference-category rates.

Measurement uses a two-pass codebook. Coders read title, description/body, OCR text, and content-bearing hashtags together; they do not infer from unobserved video, comments, author identity, or external knowledge. Pass 1 codes rhetorical mode and evidence fields; Pass 2 codes direction and carrier. Platform legibility is script-derived only after the final reviewed `message_location` field is fixed; it collapses message location, post type, visible text, OCR text, and video-observability flags into a text-observability category. Production labels are source-tracked: 623 high-risk posts use human-adjudicated final fields, while 971 non-adjudicated posts use the LLM layer only where the pre-specified rule allows it. All 623 adjudicated rows merge into the final workbook; the non-adjudicated rows have no mode consensus violations and no direction-positive leakage outside the adjudication queue.

3 Findings

RQ1: Most event participation is not a theory-bearing direction

The event corpus is not dominated by overt political claims. Public claim/correction is the largest mode (609 posts, 38.2%), but playful spectatorship is also large (441, 27.7%). Information/evidence work and inquiry/reflection account for 238 (14.9%) and 225 (14.1%) posts. Direction is rarer than event participation: the final direction dummies identify 91 anti-Manchu/Qing posts (5.7%), 14 anti-West posts (0.9%), and 98 feminist-lineage posts (6.1%). Most posts therefore participate in the Hong-Kangxi event without advancing one of the three theory-bearing directions.

Carrier and legibility explain why a direction-only design would be too coarse. Under the default carrier rule, 878 posts use Hong-Kangxi-specific material, 235 use Red Chamber suoyin material, 218 use popular-culture intertext, and 293 use broader history/politics material; these categories overlap. Legibility is also uneven: 728 posts have high caption legibility, 570 rely on medium OCR legibility, 234 have low visual/video legibility, and 54 are video-centered but still caption-readable. Appendix C reports the distributions, bivariate rates, carrier-source comparisons, and sensitivity plots.

RQ2–RQ3: The adjusted signals are feminist-lineage direction, Red Chamber suoyin carrier, and medium OCR legibility

The main model is a descriptive logit for `search_nonretrievable`. It uses the 1,463 posts with non-error outcome and codable mode/content. It includes three direction dummies, categorical mode, four carrier controls, binary `has_evidence_work`, platform legibility, post type, posting date, keyword fixed effects, log visible-text length, and log engagement at collection. Standard errors are robust. Engagement is shown both with and without controls because engagement may partly reflect earlier visibility or ranking rather than a pre-treatment covariate; Appendix 28 adds bin controls and low-engagement exclusions for this reason.

The model answers RQ2 cautiously. Feminist-lineage direction is the only theory-bearing direction that remains clearly associated with higher public-search non-retrievability in the main model and in the with/without-engagement comparison. Anti-Manchu/Qing direction is positive but not statistically distinguished from zero in the main model. Its estimate also moves across LLM-only direction variants, so it should be treated as theoretically important but empirically unresolved. Anti-West direction has only 14 positives in the full corpus and 13 with non-missing outcomes; its wide interval is low power,

Main logit with engagement

Odds ratios with robust 95% CIs. Outcome: search_nonretrievable.




Figure 1: Key adjusted associations from the main logit with engagement controls. Points are odds ratios; bars are robust 95% confidence intervals; the vertical reference line is OR = 1. Values above 1 indicate higher odds of search non-retrievability. The full coefficient plot appears in Appendix C.

not evidence of safety or absence.

For RQ3, Red Chamber suoyin carrier is the strongest non-direction association. It remains large under the GPT-default, DeepSeek, carrier-union, and carrier-intersection variants. This does not prove that Red Chamber content is moderated as literature; it shows that posts using Red Chamber as interpretive material are more often search-nonretrievable after adjusting for direction, mode, legibility, post type, keyword, date, text length, and engagement. Medium OCR legibility is also positive, though weaker after deriving legibility from the final reviewed message-location field. This points to an observability pathway: image-text or OCR-mediated posts may be indexed, verified, or searched differently from caption-primary posts. Appendix 28 shows that the core feminist-lineage and Red Chamber estimates remain positive when low-engagement posts are excluded, while the medium-OCR estimate weakens. Appendix 29 adds OCR-restricted checks. Dropping medium-OCR posts preserves the Red Chamber estimate and keeps feminist-lineage positive; the stricter high-caption-only model preserves Red Chamber but makes the feminist-lineage estimate less precise.

Bottom line for thesis design

The quantitative evidence supports a narrow claim: within this event corpus, public-search visibility friction is associated with feminist-lineage direction, Red Chamber suoyin carrier, and medium OCR legibility. The three signals likely reflect different mechanisms. Medium OCR may partly be

Table 1: Key regression results, with and without engagement controls

Term	Main model with engagement: AME OR (95% CI)		Without engagement: OR (95% CI)
Anti-West direction	0.53 [0.15, 1.81]	−0.146	0.53 [0.16, 1.73]
Anti-Manchu/Qing direction	1.27 [0.77, 2.09]	+0.057	1.29 [0.80, 2.11]
Feminist-lineage direction	2.06 [1.22, 3.46]	+0.165	1.90 [1.14, 3.16]
Hong-Kangxi carrier	0.86 [0.68, 1.09]	−0.036	0.85 [0.67, 1.08]
Red Chamber suoyin carrier	2.30 [1.64, 3.21]	+0.192	2.24 [1.61, 3.12]
Medium OCR legibility	1.37 [1.05, 1.79]	+0.073	1.33 [1.02, 1.74]
Log engagement total	0.93 [0.89, 0.97]	–	not included

Note: AME is the average discrete change reported by the notebook for binary predictors in the with-engagement model. The rare `tag_only_low_legibility` category is not interpreted because only two non-error posts fall in that cell.

search-procedure friction and low-exposure searchability; Red Chamber suoyin is robust even when low-engagement or OCR-dependent posts are removed; feminist-lineage remains positive but is more format-sensitive. The analysis does not yet support a causal claim about platform intent, nor a clean claim that anti-Manchu/Qing direction is selectively suppressed. The next stage should therefore combine repeated verification and account-level audit with qualitative close reading of the feminist-lineage and Red Chamber cases.